

Cop3402-System-Software-Final-Exam-Cheat-Sheet

Final Review System Software

With Paul Gazzillo

All Multiple Choice

Thursday, 12/04/2025, 10:00 a.m. - 12:00 p.m.

32 Multiple Choice Questions

14 Questions Midterm (Old Content)

18 Questions After Midterm (Compiler) (New Content)

12 Pages of notes - 24 Sides of Pages

Laptop with only exam open

Blank Scratch Paper and Pencil

Refresher

File Systems

① What makes the Unix file system "hierarchical"?

A hierarchical file system organizes files in a tree-like structure with a single root directory `/`.



Navigation

① What is the working directory and how do you display it?

The working directory is the directory you are currently "in" while using the terminal. It determines the starting point for relative paths.

Command to display it:

```
pwd # pwd = print working directory
```

2 What is the Unix standard command to rename a file?

```
mv old_filename.txt new_filename.txt
```

3 What Unix standard command will show you the text of a file?

```
cat filename.txt      # Display the full file  
less filename.txt    # View file page by page  
more filename.txt    # Similar to less
```

4 What does grep do?

grep searches for text patterns in files.
With no < input it will read from stdin

```
grep "hello" file.txt
```

5 How do you change the working directory to your home directory?

```
cd # No arguments will also take to home directory  
# or  
cd ~ # Takes to home directory
```

6 What is the Unix command to delete a file?

```
rm filename.txt  
rm -rf directory_path # can delete a full directory
```

7 How does the implementation of deleting a file work? Does it remove the file's contents from the storage medium?

- When you delete a file in Unix:
 1. The directory entry pointing to the file is removed.

2. The file's inode (metadata) is marked as free.
 3. The storage blocks may not be immediately overwritten.
 - The file content remains on the storage medium until reused.
 - This is why deleted files can sometimes be recovered with special tools.
-

Processes / Advanced Processes

1 How do you redirect standard output or standard input of a bash command to a file?

Redirect standard output (stdout) to a file:

```
grep "pattern" file.txt > grep.txt
```

Redirect standard input (stdin) from a file:

```
grep "pattern" < file.txt
```

Example combining input and output:

```
grep "pattern" < input.txt > grep.txt
```

2 How do you redirect standard output from one command to another command's standard input?

Use a pipe (|) to connect commands:

```
find . -type f | wc -l
```

Editor (Emacs)

1 How do you edit files in Emacs?

To open a file in Emacs, use:

```
emacs filename.txt
```

2 How do you save a file in Emacs?

While in Emacs:

C-x C-s

- Hold `Ctrl` and press `x`, then press `Ctrl` + `s`
 - Saves the current buffer to its file
-

3 How do you quit Emacs?

To quit Emacs:

C-x C-c

- Hold `Ctrl` and press `x`, then press `Ctrl` + `c`
 - Emacs will ask to save any unsaved changes before quitting
-

Version Control (Git)

1 What git command copies commits from the local repository to the remote repository?

`git push`

Sends your local commits to the remote repository (e.g., GitHub, GitLab)

Syntax: `git push <remote> <branch>`
Example: `git push origin main`

2 What git command copies commits from the remote repository to the local repository?

`git pull`

- Fetches commits from the remote repository and merges them into your local branch

- Syntax: `git pull <remote> <branch>`

Example: `git pull origin main`

3 What git command stages a new file?

```
git add filename.txt
```

- Adds the file to the **staging area** in preparation for a commit
 - To stage all files: `git add .`
-

4 What git command creates a log of the change to a staged file to the local repository?

```
git commit -m "Commit message"
```

- Commits the staged changes to the local repository
 - The `-m` flag allows you to include a short commit message describing the changes
-

File Syscalls

1 Using the open syscall to open a path given in the string `char *filepath` variable.

```
#include <fcntl.h>
#include <unistd.h>

// open file for reading only
int fd = open(filepath, O_RDONLY);
if (fd == -1) {
    perror("open"); // print error if open fails
}
```

`O_RDONLY` opens the file in read-only mode

2 How do you check for and terminate the program on an error with opening a file?

```

#include <stdlib.h>
#include <stdio.h>

int fd = open(filepath, O_RDONLY);
if (fd == -1) {
    perror("open"); // print error message
    exit(EXIT_FAILURE); // terminate program
}

```

- `perror` prints a descriptive error
 - `exit(EXIT_FAILURE)` terminates the program with a failure status
-

3 Using the `read` syscall, you already have an open file with descriptor `fd`. Read the first 200 bytes and print it to stdout.

```

#include <unistd.h>
#include <stdio.h>

char buffer[201]; // 200 bytes + null terminator
ssize_t bytesRead = read(fd, buffer, 200);

if (bytesRead == -1) {
    perror("read");
} else {
    buffer[bytesRead] = '\0'; // null-terminate string
    printf("%s", buffer); // print to stdout
}

```

- `read` returns the number of bytes actually read
 - Always check for errors (-1)
-

4 What syscall can you use to find the size and number of hard-links of a file?

```

#include <sys/stat.h>

struct stat st;
if (stat("filename.txt", &st) == 0) {
    printf("Size: %ld bytes\n", st.st_size);
    printf("Hard links: %ld\n", st.st_nlink);
}

```

`stat` retrieves metadata (size, permissions, timestamps, hard links, etc.)

5 What syscall can you use to find the name of a file?

- To read directory entries, use:

```
#include <dirent.h>

DIR *dir = opendir(".");
struct dirent *entry;
while ((entry = readdir(dir)) != NULL) {
    printf("%s\n", entry->d_name); // prints file/directory name
}
closedir(dir);
```

- `d_name` contains the filename
 - `readdir` reads one entry at a time
-

6 Print all files in a given directory, except "." and ".."

```
#include <stdio.h>
#include <dirent.h>
#include <string.h>

DIR *dir = opendir(".");
struct dirent *entry;

while ((entry = readdir(dir)) != NULL) {
    if (strcmp(entry->d_name, ".") != 0 && strcmp(entry->d_name, "..") != 0) {
        printf("%s\n", entry->d_name);
    }
}

closedir(dir);
```

- `strcmp` is used to skip `"."` and `".."` entries
 - This prints only actual files or directories in the folder
-

Process, Pipe, Syscalls

1 Write code that uses Unix standard syscalls to create a new process that runs the `ls` command

```

#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>

int main() {
    pid_t pid = fork();
    if (pid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    } else if (pid == 0) {
        // Child process
        execlp("ls", "ls", NULL);
        perror("execlp"); // only reached if exec fails
        exit(EXIT_FAILURE);
    } else {
        // Parent process waits for child
        wait(NULL);
    }
    return 0;
}

```

2 Write code that creates a new process, where the original process writes "parent" and the new process writes "child" to stdout

```

#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    pid_t pid = fork();
    if (pid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    } else if (pid == 0) {
        // Child process
        write(1, "child\n", 6); // 1 = stdout
    } else {
        // Parent process
        write(1, "parent\n", 7);
    }
    return 0;
}

```

3 Write code that replaces the current process with the `stat` or `ls` command

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    execlp("ls", "ls", NULL); // replaces current process with ls
    perror("execlp");        // only reached if exec fails
    exit(EXIT_FAILURE);
}
```

- `execlp` replaces the current process image with a new program
-

4 Write a program that opens a pipe and reads from and writes to it

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    int pipefd[2];
    char buffer[20];
    if (pipe(pipefd) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }

    pid_t pid = fork();
    if (pid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    } else if (pid == 0) {
        // Child process writes
        close(pipefd[0]);           // close read end
        write(pipefd[1], "Hello\n", 6);
        close(pipefd[1]);
    } else {
        // Parent process reads
        close(pipefd[1]);           // close write end
        read(pipefd[0], buffer, 6);
        write(1, buffer, 6);         // stdout
        close(pipefd[0]);
    }

    return 0;
}
```

5 Write a program that redirects the standard output to a file called output.txt

```
#include <unistd.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    int fd = open("output.txt", O_WRONLY | O_CREAT | O_TRUNC, 0644);
    if (fd == -1) {
        perror("open");
        exit(EXIT_FAILURE);
    }

    // Redirect stdout to file
    if (dup2(fd, 1) == -1) {
        perror("dup2");
        exit(EXIT_FAILURE);
    }

    close(fd);

    // Anything written to stdout goes to output.txt
    printf("Hello, redirected output!\n");

    return 0;
}
```

- `dup2(fd, 1)` replaces **stdout (file descriptor 1)** with the file descriptor `fd`
- All subsequent writes to `stdout` go to the file

Extras

Directories are files, but they map filenames to the `i_node` number of the file in the system.
How do you reference parent directories `<- ".."`

`open()` <- Syscall
`fopen()` <- C library call

Version control
`git pull` : Remote -> Local
`git push`: Local -> Remote

`fork()` create a new process in the unix world
`exec()` Executes a new program
`dup2()` is easy way to redirect stdout

Second Half Material

Source Code to Processes (Toolchain)

1 Know each phase of the toolchain, what it does, and their ordering

The typical C/C++ toolchain has the following phases:

| Phase | Description | Input | Output |
|------------------|---|-----------------------------------|------------------------------------|
| 1. Preprocessing | Handles <code>#include</code> , <code>#define</code> , and macros | .c or .cpp source file | Expanded source code (.i) |
| 2. Compilation | Translates preprocessed code to assembly | Preprocessed code (.i) | Assembly code (.s) |
| 3. Assembly | Converts assembly code to machine code (object code) | Assembly code (.s) | Object file (.o) |
| 4. Linking | Combines object files and libraries into an executable | Object files (.o) and libraries | Executable (a.out or custom name) |
| 5. Loading | Loads the executable into memory and starts execution | Executable file | Running process in memory |

Ordering:

Source code (.c/.cpp)

↓ Preprocessing

Preprocessed code (.i)

↓ Compilation

Assembly code (.s)

↓ Assembly

Object code (.o)

↓ Linking

Executable

↓ Loading

Running process

Preprocessing happens first, handling macros and includes

Compilation generates assembly instructions

Assembly produces machine-readable object code

Linking resolves references between files and libraries

Loading places the program into memory for execution

Compilation Pipeline

```

# = Gotten Files

# C Source (.c)

Goes through C Preprocessor (cpp)

# Preprocessed C (.i)

Goes through C Compiler (gcc)

# Assembly (.s)

Goes through Assembler (as)

# Object File (.o)

Goes through Linker (ld) -> Links 2
  > object files into 1 object file and
  > ensures the calls and callees match up

# Executable File

Goes through Loader (execve)

# Running File

```

Arithmetic

1 Understand and write both SimpleIR and assembly code for arithmetic

- **Arithmetic operations** include: addition, subtraction, multiplication, division, and modulo.
 - Both SimpleIR and assembly represent these operations, but the syntax differs.
-

2 What assembly instructions perform arithmetic?

| Operation | x86-64 Assembly Instruction |
|----------------|---|
| Addition | <code>add dest, src</code> |
| Subtraction | <code>sub dest, src</code> |
| Multiplication | <code>imul dest, src</code> |
| Division | <code>idiv src</code> (divides accumulator by src) |
| Modulo | <code>idiv src</code> (remainder stored in <code>rdx</code> after division) |

- `dest` is typically a register or memory location

- Most arithmetic instructions update processor flags for comparisons
-

3 Write equivalent assembly code

Example: Compute $c = a + b$

```
mov rax, [a]      ; load value of a into rax
add rax, [b]      ; add value of b to rax
mov [c], rax      ; store result into c
```

Example: Compute $c = a * b$

```
mov rax, [a]
imul rax, [b]
mov [c], rax
```

4 Write equivalent SimpleIR code

Example: Compute $c = a + b$

```
c := a + b
```

Example: Compute $c = a * b$

```
c := a * b
```

Bonus Stuff

```
add
sub
imul
idiv <- Integer Division
```

Be able to write assembly code

```
x = x/7
x := x / 7
```

```
Access local variables with -offset(%rbp)
ex. x -> -16(%rbp)
```

Insert Operation from Codegen

Branching

1 Understand and write both SimpleIR and assembly code for branching

- Branching changes the flow of execution based on a condition.
- Both assembly and SimpleIR can implement conditional or unconditional branches.
- Typical use cases: `if` statements, loops, and comparisons.

2 What assembly instructions perform branching?

| ASM Op | SimpleIR Op | Jump if... |
|--------|-------------|-----------------------|
| je | = | Equal |
| jne | != | Not equal |
| jl | < | Less than |
| jle | <= | Less than or equal |
| jg | > | Greater than |
| jge | >= | Greater than or equal |
| jmp | | Unconditional jump |

- Conditional jumps rely on flags set by a prior comparison (`cmp` instruction).

3 Write equivalent assembly code

Example:

```
if (a == b) {
    c = 1;
} else {
    c = 0;
}
```

Assembly (x86-64):

```
mov rax, [a]
cmp rax, [b]      ; compare a and b
je equal_label    ; jump if equal
mov rbx, 0
jmp end_label
```

```
equal_label:  
mov rbx, 1  
end_label:  
mov [c], rbx
```

4 Write equivalent SimpleIR code

Example:

```
if (a == b) { c = 1; } else { c = 0; }
```

```
t1 := a  
t2 := b  
if t1 == t2 goto true_label  
c := 0  
goto end_label  
true_label:  
c := 1  
end_label:
```

Pointers

1 Understand and write both SimpleIR and assembly code for pointers

- Pointers store the memory address of a variable.
- Assembly uses **register indirect addressing** to read/write via a pointer.
- SimpleIR can represent pointers using `load` and `store` operations with addresses.

2 Write equivalent assembly code

Example: C code:

```
int a = 5;  
int *p = &a;  
*p = 10;
```

Assembly (x86-64) using register indirect addressing:

```
mov rax, [a] ; optional: load a into rax (not needed here)  
lea rbx, [a] ; load address of a into rbx  
mov dword [rbx], 10 ; store 10 at the memory address stored in rbx
```

Explanation:

- `lea rbx, [a]` loads the address of `a` into `rbx`
 - `[rbx]` dereferences the pointer
 - `mov [rbx], 10` writes `10` into the location pointed to by `rbx`
-

3 Write equivalent SimpleIR code

Example: C code:

```
int a = 5;
int *p = &a;
*p = 10;
int x = *p;
```

Simple IR:

```
p := &a; store the address of a in p
*p := 10; store 10 at the address pointed to by p
x := *p; load the value from the address pointed to by p into x
```

Bonus Stuff

Register Copy (Load from Memory) `mov (%rax) %rbx (%rax)` <- Look up address
Parenthesis mean look up address of this register

Dereference `%rbp` at -48 offset
`mov -48(%rbp), %rax`

Working with address
`mov %rbp, %rax`
`add -48, %rax`

Register Indirect `mov`
`mov %rbx, (%rax)` <- Store the value of `rbx` into
the address of `rax`

Functions and Their Implementation (x86-64, System V ABI)

1 What are the contents of the stack frame in the x86-64 System V ABI?

A **stack frame** (activation record) typically contains:

- **Return address** – the instruction to return to after the function call
- **Saved base pointer (rbp)** – points to the previous stack frame
- **Local variables** – stored on the stack
- **Saved registers** – registers that the function must preserve according to the calling convention
- **Function parameters** (if not passed in registers)

Diagram:

```

High Memory
Function parameters (if spilled)
Saved registers
Local variables
Saved RBP
Return address
Low Memory (stack grows downward)

```

2 How are parameters passed in the x86-64 System V ABI?

- **First 6 integer/pointer arguments:** passed in registers: `rdi, rsi, rdx, rcx, r8, r9`
- **Additional arguments:** passed on the stack

Example:

```
int add(int a, int b, int c);
```

- `a` → `rdi`
- `b` → `rsi`
- `c` → `rdx`

3 How are variables represented and accessed in memory?

- **Local variables** → stored on the stack (relative to `rbp`)
- **Global/static variables** → stored in data or BSS segment
- **Pointers** → store memory addresses, accessed via dereference (`[rbp-offset]` or `[register]`)

Example:

```
mov rax, [rbp-8] ; load local variable at rbp-8 into rax
mov [rbp-16], rbx ; store rbx into local variable at rbp-16
```

4 How is the stack frame managed?

- Stack grows downward (toward lower addresses)
- Push/pop instructions or `sub rsp, <size>` allocate space for local variables
- RBP is used as a stable frame pointer to reference local variables and function parameters

Typical steps:

1. Save previous `rbp`
2. Set `rbp = rsp`
3. Allocate space for locals (`sub rsp, size`)
4. Access locals via `[rbp-offset]`

5 What do the prologue and epilogue do?

Prologue: sets up the stack frame at the start of the function

```
push rbp      ; save caller's base pointer  
mov rbp, rsp ; set base pointer for this function  
sub rsp, N   ; allocate space for local variables
```

Epilogue: cleans up the stack frame before returning

```
mov rsp, rbp ; restore stack pointer  
pop rbp      ; restore caller's base pointer  
ret          ; return to caller
```

6 How are values returned from a function in the x86-64 System V ABI?

- Integer/pointer return values: `rax`
Example:

```
mov rax, 42      ; return 42 from function  
ret
```

- Caller reads the return value from `rax`

Bonus Stuff

Stackframes are important

What information is stored in stackframe

6 Parameters in Register, the rest in stack

Some Parameters are stored in Stackframe

Caller puts Parameters in Registers and Stackframe

Return Address is on the Stackframe <- Caller Puts it there

Base Pointer [RBP or rbp] is stored on Stack <- pointing at the current active function's stackframe, Pointing at the address of the old functions base pointer.

Local variables are stored on the stack

%rip <- Current line of execution

When making a new Function call, we make a new stack frame so we always have access to the previous stack frame from the previous functional call.

How are values returned from functions

call f

- First saves the return address
- Then branches to f

ret

- ret <- Opcode pops retrieves the return address
- branches back to the return address